

```
import pandas as pd  
import numpy as np
```

Problem - 1: Perform a classification task with knn from scratch.

1. Load the Dataset:
 - Read the dataset into a pandas DataFrame.

```
df = pd.read_csv('/content/drive/MyDrive/ConceptAndTechnologiesOfAI/wi  
df.head(3)
```

	id	chol	stab.glu	hdl	ratio	glyhb	location	age	gender	he
0	1000	203.0		82	56.0	3.6	4.31	Buckingham	46	female
1	1001	165.0		97	24.0	6.9	4.44	Buckingham	29	female
2	1002	228.0		92	37.0	6.2	4.64	Buckingham	58	female

Next steps: [Generate code with df](#) [New interactive sheet](#)

- Display the first few rows and perform exploratory data analysis (EDA) to understand the dataset (e.g., check data types, missing values, summary statistics).

```
# Finding all datas  
print(df.info())  
  
print()## for spaces  
  
# Finding Missing values  
print(df.isnull().sum())  
  
print()  
  
# Finding Summary statistics  
print(df.describe(include="all"))  
  
print()  
  
# Finding Shape  
print(df.shape)  
  
print()  
  
# 6. finding names  
print(df.columns)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 403 entries, 0 to 402
Data columns (total 19 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          403 non-null    int64  
 1   chol         402 non-null    float64 
 2   stab.glu    403 non-null    int64  
 3   hdl          402 non-null    float64 
 4   ratio        402 non-null    float64 
 5   glyhb       390 non-null    float64 
 6   location     403 non-null    object  
 7   age          403 non-null    int64  
 8   gender       403 non-null    object  
 9   height       398 non-null    float64 
 10  weight       402 non-null    float64 
 11  frame         391 non-null    object  
 12  bp.1s        398 non-null    float64 
 13  bp.1d        398 non-null    float64 
 14  bp.2s        141 non-null    float64 
 15  bp.2d        141 non-null    float64 
 16  waist         401 non-null    float64 
 17  hip           401 non-null    float64 
 18  time.ppn     400 non-null    float64 
dtypes: float64(13), int64(3), object(3)
memory usage: 59.9+ KB
None
```

```
id          0
chol         1
stab.glu    0
hdl          1
ratio        1
glyhb       13
location     0
age          0
gender       0
height       5
weight       1
frame        12
bp.1s        5
bp.1d        5
bp.2s        262
bp.2d        262
waist         2
hip           2
time.ppn     3
dtype: int64
```

	id	chol	stab.glu	hdl	ratio
count	403.000000	402.000000	403.000000	402.000000	402.000000
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	15978.310174	207.845771	106.672457	50.445274	4.521642

sta	11881.122124	44.445551	25.25	220.202020	1.121880
min	1000.000000	78.000000	48.000000	12.000000	1.500000
25%	4792.500000	179.000000	81.000000	38.000000	3.200000

2. Handle Missing Data:

- Handle any missing values appropriately, either by dropping or imputing them based on the data.

```
df.isnull().sum()#find missing data
```

	0
id	0
chol	1
stab.glu	0
hdl	1
ratio	1
glyhb	13
location	0
age	0
gender	0
height	5
weight	1
frame	12
bp.1s	5
bp.1d	5
bp.2s	262
bp.2d	262
waist	2
hip	2
time.ppn	3

dtype: int64

```
df = df.drop(columns=["bp.2s", "bp.2d"]) #As there are too much empty  
df["frame"] = df["frame"].fillna(df["frame"].mode()[0]) #Fill in frame
```

```
mapping = {"small": 0, "medium": 1, "large": 2} # Adding int value
df["frame"] = df["frame"].map(mapping)

mapping = {"male": 0, "female": 1} # Adding int value
df["gender"] = df["gender"].map(mapping)

# Removing location
df = df.drop(columns=["location", "id", "chol", "stab.glu", "hdl", "glyhb"])

##Enter every other column and add the median value
num_cols = ["ratio", "height", "weight", "waist", "hip"]

for col in num_cols:
    df[col] = df[col].fillna(df[col].median())
df.head()
```

	ratio	gender	height	weight	frame	waist	hip	grid icon
0	3.6	1	62.0	121.0	1	29.0	38.0	copy icon
1	6.9	1	64.0	218.0	2	46.0	48.0	copy icon
2	6.2	1	61.0	256.0	2	49.0	57.0	
3	6.5	0	67.0	119.0	2	33.0	38.0	
4	8.9	0	68.0	183.0	1	44.0	41.0	

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
df.isnull().sum()
```

```
0  
ratio 0  
gender 0  
height 0  
weight 0  
frame 0  
waist 0  
hip 0
```

```
dtype: int64
```

3. Feature Engineering:

- Separate the feature matrix (X) and target variable (y).
- Perform a train - test split from scratch using a 70% – 30% ratio.

```
#X as everycolumns except frame  
# x = df.drop(columns=["frame"]).values  
  
important_features = ["ratio", "height", "weight", "waist", "hip"]  
  
x = df[important_features].values  
  
#Taking frame as the y  
y = df["frame"].values
```

```
def train_test_split_scratch(X, y, test_size=0.3, random_seed=42):  
    """  
        Splits dataset into train and test sets from scratch.  
    """  
  
    np.random.seed(random_seed)  
  
    # Generate array of indices  
    indices = np.arange(X.shape[0])  
    np.random.shuffle(indices)  
  
    # Determine test size  
    test_count = int(len(X) * test_size)  
  
    # Split indices  
    test_idx = indices[:test_count]  
    train_idx = indices[test_count:]  
  
    # Create train and test sets  
    X_train = X[train_idx]  
    X_test = X[test_idx]  
    y_train = y[train_idx]  
    y_test = y[test_idx]  
  
    #Printing the results  
    print("X_train shape:", X_train.shape)  
    print("X_test shape:", X_test.shape)  
    print("y_train shape:", y_train.shape)  
    print("y_test shape:", y_test.shape)  
  
    return X_train, X_test, y_train, y_test  
  
x_train, x_test, y_train, y_test = train_test_split_scratch(x, y)
```

```
X_train shape: (283, 5)
X_test shape: (120, 5)
y_train shape: (283,)
y_test shape: (120,)
```

4. Implement KNN:

- Build the KNN algorithm from scratch (no libraries like sickit-learn for KNN).
- Compute distances using Euclidean distance.
- Write functions for:
 - Predicting the class for a single query.
 - Predicting classes for all test samples.
- Evaluate the performance using accuracy.

```
def euclidean_distance(point1, point2):
    """
    Calculate the Euclidean distance between two points in n-dimensions.

    Arguments:
    point1 : np.ndarray
        The first point as a numpy array.
    point2 : np.ndarray
        The second point as a numpy array.

    Returns:
    float
        The Euclidean distance between the two points.

    Raises:
    ValueError: If the input points do not have the same dimensions.
    """
    if point1.shape != point2.shape:
        raise ValueError("Points must have the same dimensions to calculate distance.")

    distance = np.sqrt(np.sum((point1 - point2) ** 2))
    return distance

euclidean_distance(x[0], x[1])

np.float64(99.06003231000523)
```

```
try:  
  
    point1 = np.array([3, 4])  
    point2 = np.array([0, 0])  
  
    result = euclidean_distance(point1, point2)  
  
    expected_result = 5.0  
    assert np.isclose(result, expected_result), f"Expected {expected_}  
  
    print("Test passed successfully!")  
except ValueError as ve:  
    print(f"ValueError: {ve}")  
except AssertionError as ae:  
    print(f"AssertionError: {ae}")  
except Exception as e:  
    print(f"An unexpected error occurred: {e}")
```

Test passed successfully!

```
def knn_predict_single(query, X_train, y_train, k=3):  
    """  
        Predict the class label for a single query using the K-nearest ne:  
  
    Arguments:  
        query : np.ndarray  
            The query point for which the prediction is to be made.  
        X_train : np.ndarray  
            The training feature matrix.  
        y_train : np.ndarray  
            The training labels.  
        k : int, optional  
            The number of nearest neighbors to consider (default is 3).  
  
    Returns:  
        int  
            The predicted class label for the query.  
    """  
  
    distances = [euclidean_distance(query, x) for x in X_train]  
  
    sorted_indices = np.argsort(distances)  
  
    nearest_indices = sorted_indices[:k]
```

```
nearest_labels = y_train[nearest_indices]

prediction = np.bincount(nearest_labels).argmax()

return prediction

def knn_predict(X_test, X_train, y_train, k=3):
    """
    Predict the class labels for all test samples using the K-nearest

    Arguments:
    X_test : np.ndarray
        The test feature matrix.
    X_train : np.ndarray
        The training feature matrix.
    y_train : np.ndarray
        The training labels.
    k : int, optional
        The number of nearest neighbors to consider (default is 3).

    Returns:
    np.ndarray
        An array of predicted class labels for the test samples.
    """
    predictions = [knn_predict_single(x, X_train, y_train, k) for x in X_test]
    return np.array(predictions)
```

```
try:

    X_test_sample = x_test[:5]
    y_test_sample = y_test[:5]

    predictions = knn_predict(X_test_sample, x_train, y_train, k=3)

    print("Predictions:", predictions)
    print("Actual labels:", y_test_sample)

    assert predictions.shape == y_test_sample.shape, (
        "The shape of predictions does not match the shape of the actual labels"
    )

    print("Test case passed successfully!")
except AssertionError as ae:
```

```
        print(f"AssertionError: {ae}")
    except Exception as e:
        print(f"An unexpected error occurred: {e}")
```

```
Predictions: [2 1 1 2 1]
Actual labels: [1 1 1 1 2]
Test case passed successfully!
```

```
def compute_accuracy(y_true, y_pred):
    """
    Compute the accuracy of predictions.

    Arguments:
    y_true : np.ndarray
        The true labels.
    y_pred : np.ndarray
        The predicted labels.

    Returns:
    float
        The accuracy as a percentage (0 to 100).
    """
    correct_predictions = np.sum(y_true == y_pred)
    total_predictions = len(y_true)
    accuracy = (correct_predictions / total_predictions) * 100
    return accuracy
```

```
try:
    predictions = knn_predict(x_test, x_train, y_train, k=3)

    accuracy = compute_accuracy(y_test, predictions)

    print(f"Accuracy of the KNN model on the test set: {accuracy:.2f}")
except Exception as e:
    print(f"An unexpected error occurred during prediction or accuracy calculation: {e}")
```

```
Accuracy of the KNN model on the test set: 39.17%
```

```
for k in [1, 3, 5, 7, 9, 11, 13, 15, 16, 17, 18, 20]:
    preds = knn_predict(x_test, x_train, y_train, k)
    acc = np.mean(preds == y_test)
    print(f"k = {k} -> Accuracy = {acc*100:.2f}%")
```

```
k = 1 -> Accuracy = 43.33%
k = 3 -> Accuracy = 39.17%
k = 5 -> Accuracy = 49.17%
k = 7 -> Accuracy = 45.83%
k = 9 -> Accuracy = 46.67%
k = 11 -> Accuracy = 47.50%
k = 13 -> Accuracy = 48.33%
k = 15 -> Accuracy = 48.33%
k = 16 -> Accuracy = 52.50%
k = 17 -> Accuracy = 50.83%
k = 18 -> Accuracy = 49.17%
k = 20 -> Accuracy = 49.17%
```

```
import matplotlib.pyplot as plt
def experiment_knn_k_values(X_train, y_train, X_test, y_test, k_values):
    """
        Run KNN predictions for different values of k and plot the accuracy.

    Arguments:
        X_train : np.ndarray
            The training feature matrix.
        y_train : np.ndarray
            The training labels.
        X_test : np.ndarray
            The test feature matrix.
        y_test : np.ndarray
            The test labels.
        k_values : list of int
            A list of k values to experiment with.

    Returns:
        dict
            A dictionary with k values as keys and their corresponding accuracy as values.
    """
    accuracies = {}

    for k in k_values:

        predictions = knn_predict(X_test, X_train, y_train, k=k)

        accuracy = compute_accuracy(y_test, predictions)
        accuracies[k] = accuracy

        print(f"Accuracy for k={k}: {accuracy:.2f}%")

    plt.figure(figsize=(10, 5))
    plt.plot(k_values, list(accuracies.values()), marker='o')
    plt.xlabel('k (Number of Neighbors)')
```

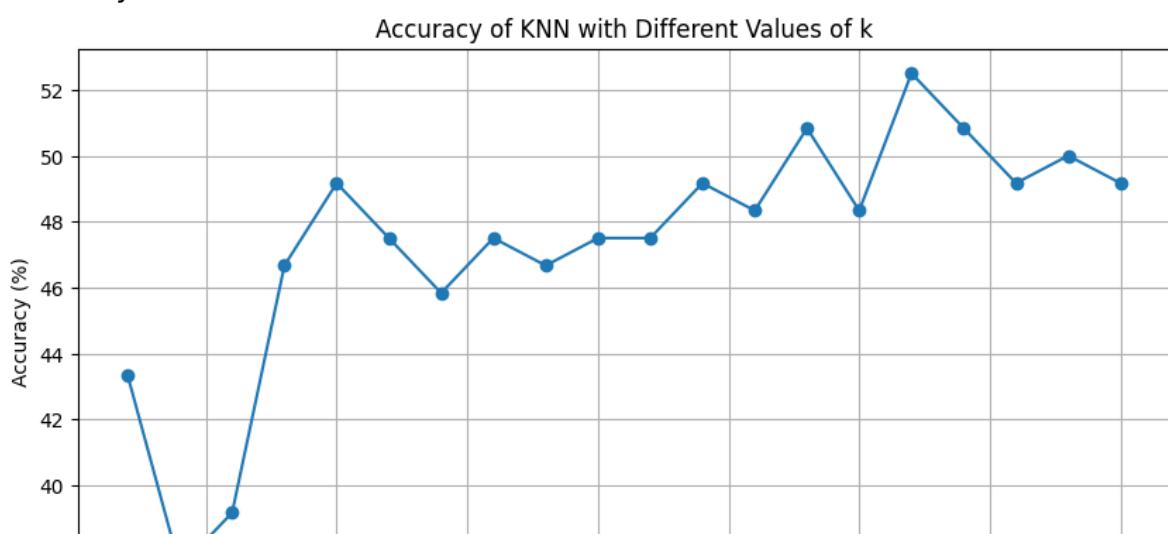
```
plt.ylabel('Accuracy (%)')
plt.title('Accuracy of KNN with Different Values of k')
plt.grid(True)
plt.show()

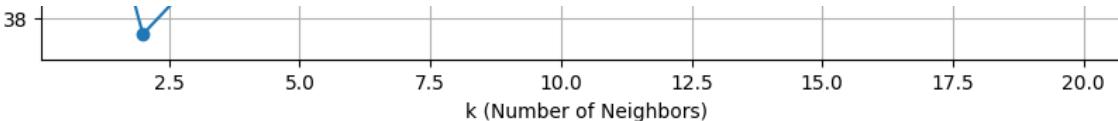
return accuracies

k_values = range(1, 21)

try:
    accuracies = experiment_knn_k_values(x_train, y_train, x_test, y_
    print("Experiment completed. Check the plot for the accuracy trend")
except Exception as e:
    print(f"An unexpected error occurred during the experiment: {e}")
```

```
Accuracy for k=1: 43.33%
Accuracy for k=2: 37.50%
Accuracy for k=3: 39.17%
Accuracy for k=4: 46.67%
Accuracy for k=5: 49.17%
Accuracy for k=6: 47.50%
Accuracy for k=7: 45.83%
Accuracy for k=8: 47.50%
Accuracy for k=9: 46.67%
Accuracy for k=10: 47.50%
Accuracy for k=11: 47.50%
Accuracy for k=12: 49.17%
Accuracy for k=13: 48.33%
Accuracy for k=14: 50.83%
Accuracy for k=15: 48.33%
Accuracy for k=16: 52.50%
Accuracy for k=17: 50.83%
Accuracy for k=18: 49.17%
Accuracy for k=19: 50.00%
Accuracy for k=20: 49.17%
```





Experiment completed. Check the plot for the accuracy trend.

Problem - 2 - Experimentation:

1. Repeat the Classification Task:
 - Scale the Feature matrix X.
 - Use the scaled data for training and testing the kNN Classifier.
 - Record the results.
2. Comparative Analysis: Compare the Results -
 - Compare the accuracy and performance of the kNN model on the original dataset from problem 1 versus the scaled dataset.
 - Discuss:
 - How scaling impacted the KNN performance.
 - The reason for any observed changes in accuracy.

```
def min_max_scale(X):  
    """  
        Manually scale the feature matrix X using min-max scaling.  
        Returns the scaled version of X.  
    """  
    X_min = X.min(axis=0)      # minimum of each column  
    X_max = X.max(axis=0)      # maximum of each column  
  
    return (X - X_min) / (X_max - X_min)  
  
x_scaled = min_max_scale(x)  
x_scaled  
  
array([[0.11797753, 0.41666667, 0.09734513, 0.1       , 0.23529412],  
       [0.3033708 , 0.5       , 0.52654867, 0.66666667, 0.52941176],  
       [0.26404494, 0.375     , 0.69469027, 0.76666667, 0.79411765],  
       ...,  
       [0.20224719, 0.70833333, 0.30088496, 0.23333333, 0.29411765],  
       [0.12921349, 0.45833333, 0.43362832, 0.5       , 0.52941176],  
       [0.02808080, 0.5       , 0.52520273, 0.76666667, 0.822520411111]
```

```
↳ .v.02000909, v.0 , v.00000000, v.02002741]]
```

```
x_train_s, x_test_s, y_train_s, y_test_s = train_test_split_scratch(x,
```

```
X_train shape: (283, 5)
X_test shape: (120, 5)
y_train shape: (283,)
y_test shape: (120,)
```

```
y_pred_scaled = knn_predict(x_test_s, x_train_s, y_train_s, k=3)
y_pred_scaled
```

```
array([0, 1, 1, 2, 1, 1, 0, 0, 1, 1, 2, 1, 2, 2, 2, 0, 1, 1, 1, 0, 1,
       0, 0, 2, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 2, 1,
       1, 0, 0, 0, 1, 2, 2, 1, 1, 1, 0, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1,
       2, 1, 0, 0, 0, 2, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 2, 1,
       1, 0, 0, 1, 1, 0, 2, 1, 1, 1, 0, 1, 0, 0, 1, 2, 1, 1, 0, 1, 1,
       0, 1, 0, 2, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1])
```

```
accuracy_scaled = np.mean(y_pred_scaled == y_test_s)
accuracy_scaled
```

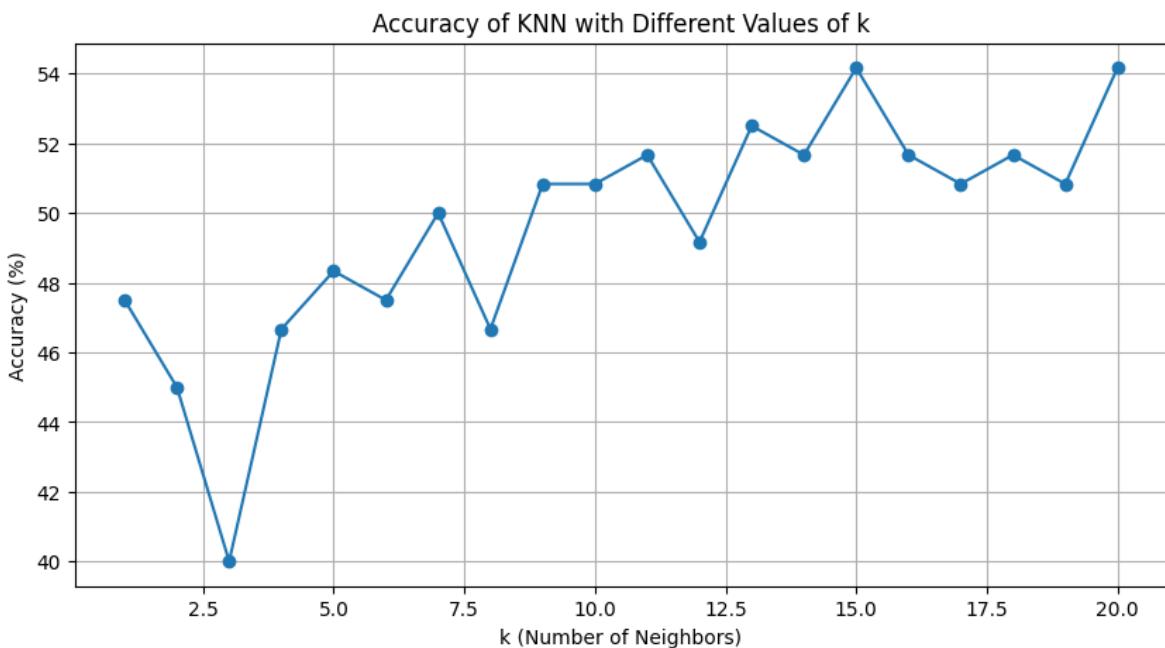
```
np.float64(0.4)
```

```
k_values = range(1, 21)
```

```
try:
    accuracies = experiment_knn_k_values(x_train_s, y_train_s, x_test_s)
    print("Experiment completed. Check the plot for the accuracy trend")
except Exception as e:
    print(f"An unexpected error occurred during the experiment: {e}")
```

```
Accuracy for k=1: 47.50%
Accuracy for k=2: 45.00%
Accuracy for k=3: 40.00%
Accuracy for k=4: 46.67%
Accuracy for k=5: 48.33%
Accuracy for k=6: 47.50%
Accuracy for k=7: 50.00%
Accuracy for k=8: 46.67%
Accuracy for k=9: 50.83%
Accuracy for k=10: 50.83%
Accuracy for k=11: 51.67%
Accuracy for k=12: 49.17%
```

```
Accuracy for k=12: 54.17%
Accuracy for k=13: 52.50%
Accuracy for k=14: 51.67%
Accuracy for k=15: 54.17%
Accuracy for k=16: 51.67%
Accuracy for k=17: 50.83%
Accuracy for k=18: 51.67%
Accuracy for k=19: 50.83%
Accuracy for k=20: 54.17%
```



Experiment completed. Check the plot for the accuracy trend.

```
##Scaling improved the performance of the kNN model.  
# Since kNN uses Euclidean distance, features with large values dominate.
```

Before scaling there were large valued features and the smaller features had little effect. And lower accuracy was observed

After Scaling all features were normalized to the same scale so the accuracy improved

Problem - 3 - Experimentation with k:

1. Vary the number of neighbors - k:
- Run the KNN model on both the original and scaled datasets for a range of:

k= 1, 2, 3, . . . 15

- For each k, record:
 - Accuracy.
 - Time taken to make predictions.

2. Visualize the Results:

- Plot the following graphs:
 - k vs. Accuracy for original and scaled datasets.
 - k vs. Time Taken for original and scaled datasets.

3. Analyze and Discuss:

- Discuss how the choice of k affects the accuracy and computational cost.
- Identify the optimal k based on your analysis.

```
import time

k_values = range(1, 16)

scaled_accuracy = []
unscaled_accuracy = []
time_unscaled = []
time_scaled = []

for kk in k_values:
    # Unscaled
    start = time.time()
    unscaled_pred = knn_predict(x_test, x_train, y_train, kk)
    end = time.time()

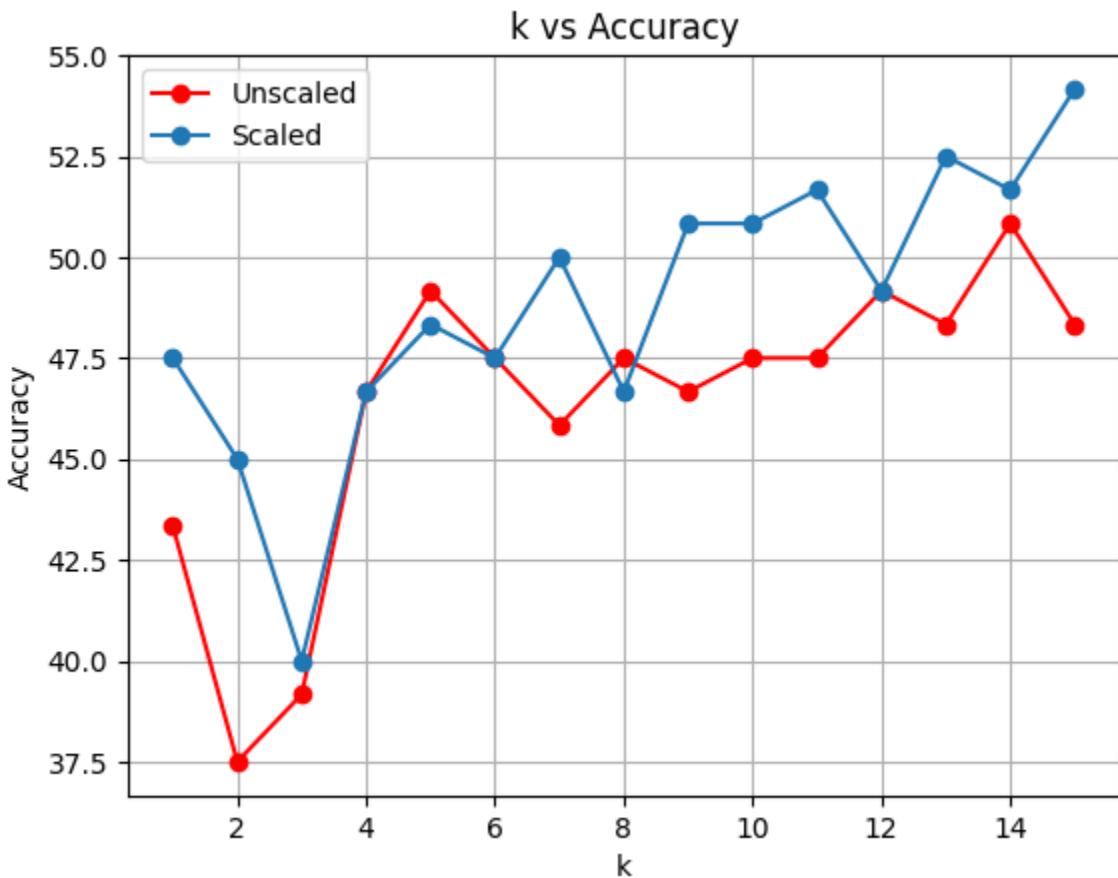
    unscaled_accuracy.append(compute_accuracy(y_test, unscaled_pred))
    time_unscaled.append(end - start)

    # Scaled
    start = time.time()
    scaled_pred = knn_predict(x_test_s, x_train_s, y_train, kk)
    end = time.time()

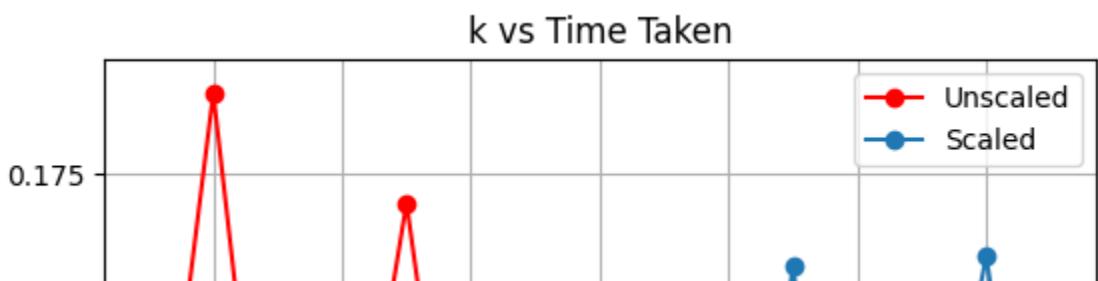
    scaled_accuracy.append(compute_accuracy(y_test, scaled_pred))
    time_scaled.append(end - start)
```

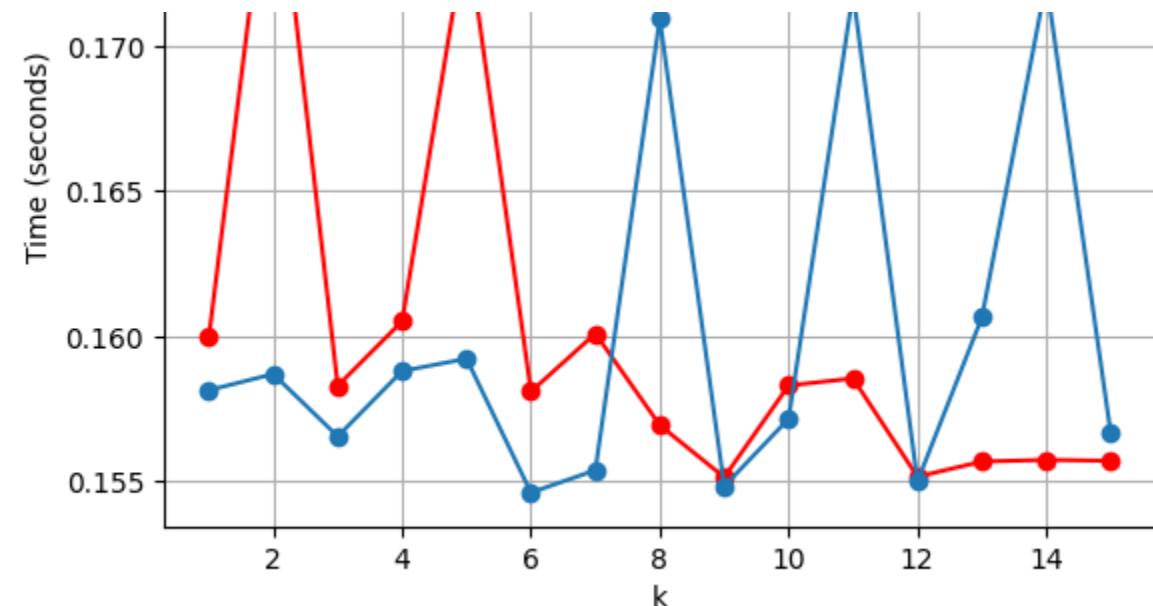
```
plt.plot(k_values, unscaled_accuracy, marker='o', label='Unscaled', col
nlt plt.plot(k values scaled accuracy marker='o' label='Scaled')
```

```
plt.plot(k_values, scaled_accuracy, marker='o', label='Scaled',  
         plt.xlabel('k')  
         plt.ylabel('Accuracy')  
         plt.title('k vs Accuracy')  
         plt.legend()  
         plt.grid(True)  
         plt.show()
```



```
plt.plot(k_values, time_unscaled, marker='o', label='Unscaled', color='red')  
plt.plot(k_values, time_scaled, marker='o', label='Scaled')  
plt.xlabel('k')  
plt.ylabel('Time (seconds)')  
plt.title('k vs Time Taken')  
plt.legend()  
plt.grid(True)  
plt.show()
```





The accuracy of the KNN model is low and unstable for small values of k because the model is sensitive to noise. As k increases, accuracy improves since predictions are based on more neighbors, making them more stable. The scaled dataset consistently performs better than the unscaled dataset because scaling ensures fair distance calculations. The prediction time remains almost constant for all values of k, so increasing k does not significantly increase computational cost. Based on the plots, the optimal value of k lies around 13–15, where the scaled data achieves the highest accuracy with reasonable computation time.

Problem - 4 - Additional Questions {Optional - But Highly Recommended}:

- Discuss the challenges of using KNN for large datasets and high-dimensional data.
- Suggest strategies to improve the efficiency of KNN (e.g., approximate nearest neighbors, dimensionality reduction).

T B I <> ⊞ ☒ “ ≡ ≡ – ψ ☺

Close

Challenges of using KNN for large datasets and high-dimensional data: KNN becomes computationally expensive for large datasets because it calculates distances between every test sample and all training samples. It also consumes a lot of memory as it stores all training data. For high-dimensional data, the "curse of dimensionality" makes

Strategies to improve KNN efficiency: Efficiency can be improved by using a KD-Trees or Ball Trees, which reduce

Challenges of using KNN for large datasets and high-dimensional data: KNN becomes computationally expensive for large datasets because it calculates distances between every test sample and all training samples. It also consumes a lot of memory as it stores all training data. For high-dimensional data, the "curse of dimensionality" makes

Dimensionality reduction techniques simplify the feature space, making distances more meaningful. This ensures no single feature dominates the distance metric. Selecting the optimal K through cross-validation improves accuracy. Clustering can also be used to limit search space for faster predictions.

The curse of dimensionality makes distances less meaningful, causing predictions to become less accurate. Additionally, KNN is sensitive to irrelevant features and feature scales, which can further reduce its performance.

Strategies to improve KNN efficiency:
Efficiency can be improved by using approximate nearest neighbor methods like KD-Trees or Ball Trees, which reduce the number of distance calculations. Dimensionality reduction techniques such as PCA can help simplify the feature space, making distances more meaningful and faster to compute. Feature scaling ensures no single feature dominates the distance metric, and selecting the optimal K through cross-validation improves accuracy. Clustering can also be used to limit search space for faster predictions.

